

# Logistic Regression

Prajwal Vijaywargiya - 2017B3A70954H

Siddhi Mahesh Burse - 2017B3A70972H

Parth Krishna Sharma - 2017B3A70907H

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

```
In [2]: train = pd.read_csv("D:/BITS/DL/Assignments/Assignment1/mnist_train.csv")
test = pd.read_csv("D:/BITS/DL/Assignments/Assignment1/mnist_test.csv")
```

```
In [3]: X_train = train.iloc[:,1:].values
y_train = train.iloc[:,0].values
```

```
In [4]: X_test = test.iloc[:,1:].values
y_test = test.iloc[:,0].values
```

```
In [5]: X_train,X_test = X_train/255.0,X_test/255.0
```

```
In [6]: lr = LogisticRegression(max_iter=1000,verbose=2)
lr.fit(X_train,y_train)
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 3.4min remaining: 0.0s
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 3.4min finished
```

```
Out[6]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
        intercept_scaling=1, l1_ratio=None, max_iter=1000,
        multi_class='auto', n_jobs=None, penalty='l2',
        random_state=None, solver='lbfgs', tol=0.0001, verbose=2,
        warm_start=False)
```


```
In [7]: pred = lr.predict(X_test)
```

```
In [8]: acc = accuracy_score(y_test,pred)
acc
```

```
Out[8]: 0.9256
```

```
In [9]: sns.heatmap(confusion_matrix(y_test,pred),cmap='Blues_r',fmt='',annot=True,cbar=False)
plt.ylabel("Correct Label")
plt.xlabel("Predicted Label")
plt.title("Confusion Matrix for Logistic Regression\nModel Accuracy={}%"
.format(acc*100))
```

```
Out[9]: Text(0.5, 1.0, 'Confusion Matrix for Logistic Regression\nModel Accuracy=92.56%')
```



```
In [11]: print("Classification Report for Logistic Regression:\n{}".format(classification_report(y_test,pred)))
```

```
Classification Report for Logistic Regression:
              precision    recall  f1-score   support

    0           0.95         0.97         0.96         980
    1           0.96         0.98         0.97        1135
    2           0.93         0.90         0.91        1032
    3           0.90         0.92         0.91        1010
    4           0.94         0.94         0.94         982
    5           0.90         0.87         0.88         892
    6           0.94         0.95         0.95         958
    7           0.93         0.92         0.93        1028
    8           0.88         0.88         0.88         974
    9           0.91         0.92         0.91        1009

 accuracy              0.93        10000
 macro avg           0.92         0.92         0.92        10000
 weighted avg        0.93         0.93         0.93        10000
```

```
In [ ]:
```